

BHOOMADI LIKHITHA REDDY PRODIGY TASK 3

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: import warnings

warnings.filterwarnings('ignore')
```

```
In [5]: data = "C:\\Users\\DELL\\Downloads\\car_evaluation.csv"
df = pd.read_csv(data, header=None)
```

```
In [7]: df
```

```
Out[7]:
```

	0	1	2	3	4	5	6
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc
...
1723	low	low	5more	more	med	med	good
1724	low	low	5more	more	med	high	vgood
1725	low	low	5more	more	big	low	unacc
1726	low	low	5more	more	big	med	good
1727	low	low	5more	more	big	high	vgood

1728 rows × 7 columns

```
In [9]: df.shape
```

```
Out[9]: (1728, 7)
```

```
In [11]: df.head()
```

```
Out[11]:
```

	0	1	2	3	4	5	6
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
In [13]: col_names = ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']

df.columns = col_names

col_names
```

```
Out[13]: ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
```

```
In [15]: df.head()
```

```
Out[15]:
```

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
In [17]: df.info
```

```
Out[17]: <bound method DataFrame.info of
ass
0    vhigh  vhigh    2    2  small  low  unacc
1    vhigh  vhigh    2    2  small  med  unacc
2    vhigh  vhigh    2    2  small  high unacc
3    vhigh  vhigh    2    2    med  low  unacc
4    vhigh  vhigh    2    2    med  med  unacc
...     ...     ...     ...     ...     ...     ...     ...
1723   low    low  5more  more    med  med  good
1724   low    low  5more  more    med  high vgood
1725   low    low  5more  more    big  low  unacc
1726   low    low  5more  more    big  med  good
1727   low    low  5more  more    big  high vgood

[1728 rows x 7 columns]>
```

```
In [19]: col_names = ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']

for col in col_names:

    print(df[col].value_counts())
```

```

buying
vhigh    432
high     432
med      432
low      432
Name: count, dtype: int64
maint
vhigh    432
high     432
med      432
low      432
Name: count, dtype: int64
doors
2         432
3         432
4         432
5more     432
Name: count, dtype: int64
persons
2         576
4         576
more      576
Name: count, dtype: int64
lug_boot
small     576
med       576
big       576
Name: count, dtype: int64
safety
low       576
med       576
high      576
Name: count, dtype: int64
class
unacc    1210
acc       384
good       69
vgood     65
Name: count, dtype: int64

```

```
In [21]: df['class'].value_counts()
```

```

Out[21]: class
unacc    1210
acc       384
good       69
vgood     65
Name: count, dtype: int64

```

```
In [23]: df.isnull().sum()
```

```

Out[23]: buying      0
maint      0
doors      0
persons    0
lug_boot   0
safety     0
class      0
dtype: int64

```

```
In [25]: # Declare feature vector and target variable
X = df.drop(['class'], axis=1)

y = df['class']
```

```
In [27]: # split X and y into training and testing sets

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.33, random_sta
```

```
In [29]: # check the shape of X_train and X_test

X_train.shape, X_test.shape
```

```
Out[29]: ((1157, 6), (571, 6))
```

```
In [31]: #Feature Engineering
```

```
In [33]: # check data types in X_train

X_train.dtypes
```

```
Out[33]: buying      object
maint      object
doors      object
persons    object
lug_boot   object
safety     object
dtype: object
```

```
In [35]: X_train.head()
```

```
Out[35]:
```

	buying	maint	doors	persons	lug_boot	safety
48	vhigh	vhigh	3	more	med	low
468	high	vhigh	3	4	small	low
155	vhigh	high	3	more	small	high
1721	low	low	5more	more	small	high
1208	med	low	2	more	small	high

```
In [37]: !pip install category_encoders
```

Collecting category_encoders

Obtaining dependency information for category_encoders from https://files.pythonhosted.org/packages/7f/e5/79a62e5c9c9ddbfa9ff5222240d408c1eeea4e38741a0dc8343edc7ef1ec/c/category_encoders-2.6.3-py2.py3-none-any.whl.metadata

Downloading category_encoders-2.6.3-py2.py3-none-any.whl.metadata (8.0 kB)

Requirement already satisfied: numpy>=1.14.0 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (1.24.3)

Requirement already satisfied: scikit-learn>=0.20.0 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (1.3.0)

Requirement already satisfied: scipy>=1.0.0 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (1.11.1)

Requirement already satisfied: statsmodels>=0.9.0 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (0.14.0)

Requirement already satisfied: pandas>=1.0.5 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (2.0.3)

Requirement already satisfied: patsy>=0.5.1 in c:\users\dell\anaconda3\lib\site-packages (from category_encoders) (0.5.3)

Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\dell\anaconda3\lib\site-packages (from pandas>=1.0.5->category_encoders) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\dell\anaconda3\lib\site-packages (from pandas>=1.0.5->category_encoders) (2023.3.post1)

Requirement already satisfied: tzdata>=2022.1 in c:\users\dell\anaconda3\lib\site-packages (from pandas>=1.0.5->category_encoders) (2023.3)

Requirement already satisfied: six in c:\users\dell\anaconda3\lib\site-packages (from patsy>=0.5.1->category_encoders) (1.16.0)

Requirement already satisfied: joblib>=1.1.1 in c:\users\dell\anaconda3\lib\site-packages (from scikit-learn>=0.20.0->category_encoders) (1.2.0)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\dell\anaconda3\lib\site-packages (from scikit-learn>=0.20.0->category_encoders) (2.2.0)

Requirement already satisfied: packaging>=21.3 in c:\users\dell\anaconda3\lib\site-packages (from statsmodels>=0.9.0->category_encoders) (23.1)

Downloading category_encoders-2.6.3-py2.py3-none-any.whl (81 kB)

----- 0.0/81.9 kB ? eta -:--:--

----- 0.0/81.9 kB ? eta -:--:--

----- 10.2/81.9 kB ? eta -:--:--

----- 61.4/81.9 kB 656.4 kB/s eta 0:00:01

----- 81.9/81.9 kB 761.9 kB/s eta 0:00:00

Installing collected packages: category_encoders

Successfully installed category_encoders-2.6.3

In [39]: `# import category encoders`

```
import category_encoders as ce
```

In [41]: `# encode variables with ordinal encoding`

```
encoder = ce.OrdinalEncoder(cols=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'price'])
```

```
X_train = encoder.fit_transform(X_train)
```

```
X_test = encoder.transform(X_test)
```

```
X_train.head()
```

Out[41]:

	buying	maint	doors	persons	lug_boot	safety
48	1	1	1	1	1	1
468	2	1	1	2	2	1
155	1	2	1	1	2	2
1721	3	3	2	1	2	2
1208	4	3	3	1	2	2

In [43]: `X_test.head()`

Out[43]:

	buying	maint	doors	persons	lug_boot	safety
599	2	2	4	3	1	2
1201	4	3	3	2	1	3
628	2	2	2	3	3	3
1498	3	2	2	2	1	3
1263	4	3	4	1	1	1

In [45]: *#Decision Tree Classifier with criterion gini index*

In [47]: *# import DecisionTreeClassifier*

```
from sklearn.tree import DecisionTreeClassifier
```

In [49]: *# instantiate the DecisionTreeClassifier model with criterion gini index*

```
clf_gini = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=0)
```

fit the model

```
clf_gini.fit(X_train, y_train)
```

Out[49]:

▼
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=3, random_state=0)

In [51]: *#Predict the Test set results with criterion gini index*

In [53]: `y_pred_gini = clf_gini.predict(X_test)`

In [55]: *#Check accuracy score with criterion gini index*

In [57]: `from sklearn.metrics import accuracy_score`

```
print('Model accuracy score with criterion gini index: {0:0.4f}'.format(accuracy_score(y_test, y_pred_gini)))
```

Model accuracy score with criterion gini index: 0.8021

```
In [59]: y_pred_train_gini = clf_gini.predict(X_train)
```

```
y_pred_train_gini
```

```
Out[59]: array(['unacc', 'unacc', 'unacc', ..., 'unacc', 'unacc', 'acc'],
      dtype=object)
```

```
In [61]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_t
```

```
Training-set accuracy score: 0.7865
```

```
In [63]: #Check for overfitting and underfitting
# print the scores on training and test set
```

```
print('Training set score: {:.4f}'.format(clf_gini.score(X_train, y_train)))
```

```
print('Test set score: {:.4f}'.format(clf_gini.score(X_test, y_test)))
```

```
Training set score: 0.7865
```

```
Test set score: 0.8021
```


```
In [65]: #Visualize decision-trees
```

```
In [67]: plt.figure(figsize=(12,8))
```

```
from sklearn import tree
```

```
tree.plot_tree(clf_gini.fit(X_train, y_train))
```

```
Out[67]: [Text(0.4, 0.875, 'x[5] <= 1.5\ngini = 0.455\nsamples = 1157\nvalue = [255, 49, 813,
40]'),
  Text(0.2, 0.625, 'gini = 0.0\nsamples = 386\nvalue = [0, 0, 386, 0]'),
  Text(0.6, 0.625, 'x[3] <= 2.5\ngini = 0.577\nsamples = 771\nvalue = [255, 49, 427, 4
0]'),
  Text(0.4, 0.375, 'x[0] <= 2.5\ngini = 0.631\nsamples = 525\nvalue = [255, 49, 181, 4
0]'),
  Text(0.2, 0.125, 'gini = 0.496\nsamples = 271\nvalue = [124, 0, 147, 0]'),
  Text(0.6, 0.125, 'gini = 0.654\nsamples = 254\nvalue = [131, 49, 34, 40]'),
  Text(0.8, 0.375, 'gini = 0.0\nsamples = 246\nvalue = [0, 0, 246, 0])]
```



```

In [69]: # Decision Tree Classifier with criterion entropy
# instantiate the DecisionTreeClassifier model with criterion entropy

clf_en = DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=0)

# fit the model
clf_en.fit(X_train, y_train)

```

```

Out[69]: ▼ DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=0)

```

```

In [71]: #Predict the Test set results with criterion entropy
y_pred_en = clf_en.predict(X_test)
#Check accuracy score with criterion entropy
from sklearn.metrics import accuracy_score

print('Model accuracy score with criterion entropy: {0:0.4f}'.format(accuracy_score(y
Model accuracy score with criterion entropy: 0.8021

```

```

In [73]: #Compare the train-set and test-set accuracy

```

```

In [75]: y_pred_train_en = clf_en.predict(X_train)

```

```

In [77]: y_pred_train_en

```

```

Out[77]: array(['unacc', 'unacc', 'unacc', ..., 'unacc', 'unacc', 'acc'],
dtype=object)

```



```
In [79]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_t
Training-set accuracy score: 0.7865

In [81]: #Check for overfitting and underfitting

In [83]: # print the scores on training and test set

print('Training set score: {:.4f}'.format(clf_en.score(X_train, y_train)))

print('Test set score: {:.4f}'.format(clf_en.score(X_test, y_test)))

Training set score: 0.7865
Test set score: 0.8021
```

We can see that the training-set score and test-set score is same as above.

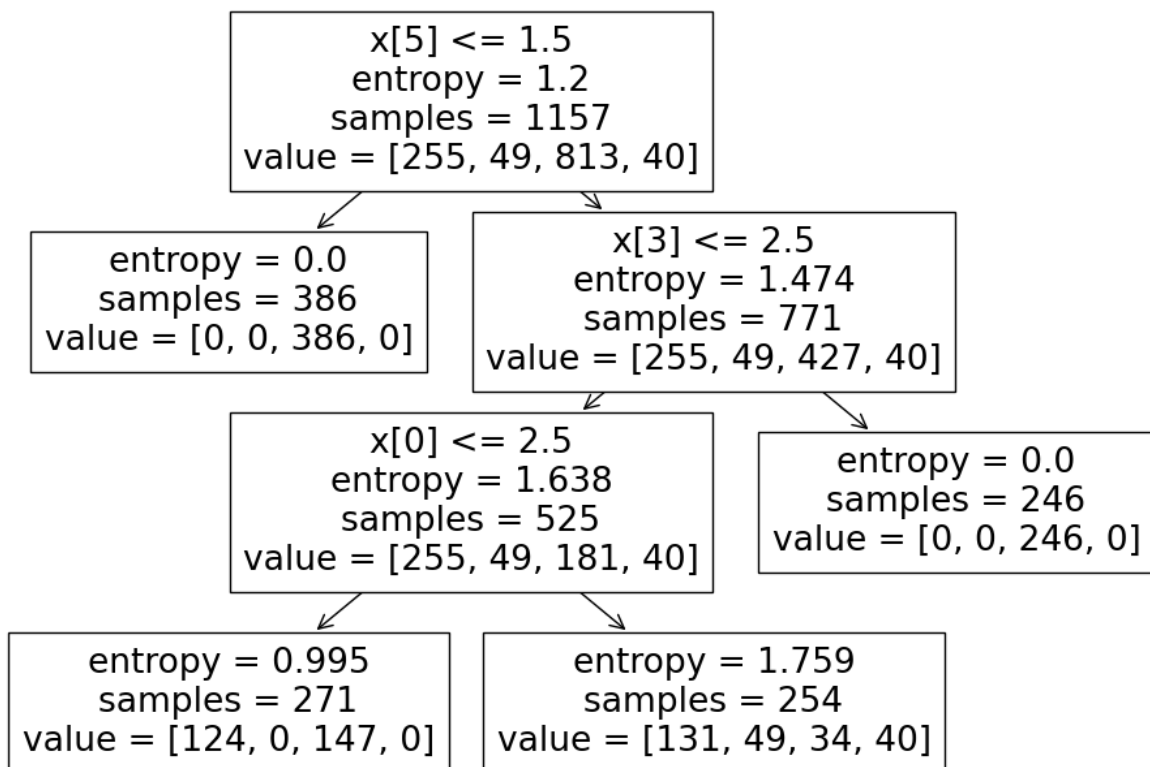
The training-set accuracy score is 0.7865 while the test-set accuracy to be 0.8021. These two values are quite comparable. So, there is no sign of overfitting.

```
In [87]: #Visualize decision-trees
plt.figure(figsize=(12,8))

from sklearn import tree

tree.plot_tree(clf_en.fit(X_train, y_train))

Out[87]: [Text(0.4, 0.875, 'x[5] <= 1.5\nentropy = 1.2\nsamples = 1157\nvalue = [255, 49, 813,
40]'),
Text(0.2, 0.625, 'entropy = 0.0\nsamples = 386\nvalue = [0, 0, 386, 0]'),
Text(0.6, 0.625, 'x[3] <= 2.5\nentropy = 1.474\nsamples = 771\nvalue = [255, 49, 42
7, 40]'),
Text(0.4, 0.375, 'x[0] <= 2.5\nentropy = 1.638\nsamples = 525\nvalue = [255, 49, 18
1, 40]'),
Text(0.2, 0.125, 'entropy = 0.995\nsamples = 271\nvalue = [124, 0, 147, 0]'),
Text(0.6, 0.125, 'entropy = 1.759\nsamples = 254\nvalue = [131, 49, 34, 40]'),
Text(0.8, 0.375, 'entropy = 0.0\nsamples = 246\nvalue = [0, 0, 246, 0]')]
```



Now, based on the above analysis we can conclude that our classification model accuracy is very good.

Our model is doing a very good job in terms of predicting the class labels.

But, it does not give the underlying distribution of values. Also, it does not tell anything about the type of errors our classifier is making.

We have another tool called Confusion matrix that comes to our rescue

```
In [90]: # Confusion matrix
```

```
In [92]: # Print the Confusion Matrix and slice it into four pieces
```

```
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred_en)

print('Confusion matrix\n\n', cm)
```

Confusion matrix

```
[[ 73   0  56   0]
 [ 20   0   0   0]
 [ 12   0 385   0]
 [ 25   0   0   0]]
```

In [94]: `#Classification Report`

In [96]: `from sklearn.metrics import classification_report`
`print(classification_report(y_test, y_pred_en))`

	precision	recall	f1-score	support
acc	0.56	0.57	0.56	129
good	0.00	0.00	0.00	20
unacc	0.87	0.97	0.92	397
vgood	0.00	0.00	0.00	25
accuracy			0.80	571
macro avg	0.36	0.38	0.37	571
weighted avg	0.73	0.80	0.77	571

Results and conclusion

1. In this project, I build a Decision-Tree Classifier model to predict the safety of the car.

I build two models, one with criterion gini index and another one with criterion entropy. The model yields a very good performance as indicated by the model accuracy in both the cases which was found to be 0.8021.

1. In the model with criterion gini index, the training-set accuracy score is 0.7865 while the test-set accuracy to be 0.8021. These two values are quite comparable. So, there is no sign of overfitting.
3. Similarly, in the model with criterion entropy, the training-set accuracy score is 0.7865 while the test-set accuracy to be 0.8021. We get the same values as in the case with criterion gini. So, there is no sign of overfitting.
4. In both the cases, the training-set and test-set accuracy score is the same. It may happen because of small dataset.
5. The confusion matrix and classification report yields very good model performance.

In []: